

# Spanner: Google's Globally-Distributed Database\*

Kocher Daniel

Summarizing Talk  
Efficient Algorithms Seminar  
University of Salzburg

January 29, 2016

\* James C. Corbett et al. (26 authors), OSDI 2012.

# Outline

- 1 Motivation
- 2 Implementation
- 3 TrueTime
- 4 Concurrency Control
- 5 Evaluation
- 6 Conclusions & Future Work

# Motivation

- Distributed data at global scale
- Externally-consistent distributed transactions
- High availability

- Distributed data at global scale
- Externally-consistent distributed transactions
- High availability
- Why not use Bigtable?
  - Difficult to use for applications with complex, evolving schemas
  - Only **eventually-consistent** (no strong consistency)

- Distributed data at global scale
- Externally-consistent distributed transactions
- High availability
- Why not use Bigtable?
  - Difficult to use for applications with complex, evolving schemas
  - Only **eventually-consistent** (no strong consistency)
- Why not use Megastore?
  - Poor performance

**Spanner** evolved from a Bigtable-like versioned key-value store to a **temporal multi-version database**

**Spanner** evolved from a Bigtable-like versioned key-value store to a **temporal multi-version database**

- Versioned data is stored in schematized semi-relational tables
- Each version is automatically timestamped with its commit time
- Garbage collection & reads at old timestamps
- **General-purpose transactions**
- SQL-like query language

- **External consistency:**

If  $T_1$  preceeds  $T_2$ , then  $T_1$  is serialized first

- **Paxos (Replication):**

- Solves problem of resilient replication of data
- Data eventually propagates to all nodes
- Different nodes always see the same data
- Majority of nodes up  $\Rightarrow$  Writes/Reads processed correctly
- Single node is elected as leader and initiates consensus
- Guarantees **consistency**



- **Two-Phase Locking (2PL):**

Guarantees serializability

- ① *Expanding* phase: locks are acquired, no locks are released
- ② *Shrinking* phase: locks are released, no locks are acquired

- **Two-Phase Commit (2PC):**

Coordinates processes participating in atomic distributed transaction

- ① *Commit-Request* phase: Request "Yes" (Commit) or "No" (Abort) from every transaction process
- ② *Commit* phase: Commit transaction if all voted "Yes", otherwise abort

As a globally-distributed database, Spanner provides interesting features:

- Dynamically controlled replication configurations at fine grain
- Externally consistent reads & writes
- globally-consistent reads at a timestamp

⇒ enables **atomic schema changes** in presence of ongoing transactions

- Timestamps reflect serialization order
- Key enabler: novel **TrueTime API exposing clock uncertainty**

# Implementation

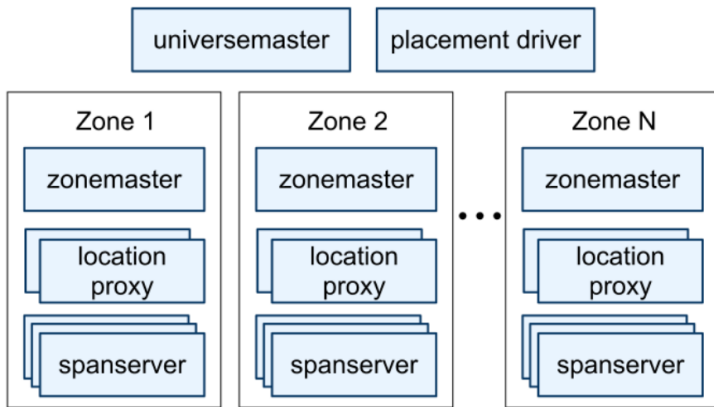


Figure 1 : Spanner server organization

- Responsible for 100 - 1000 **tablets**
- A tablet is a bag of key-value mappings  
(key:string, timestamp:int64) → string
- Timestamps are assigned to data ⇒ **multi-version database**
- Tablet states are stored on Colossus
- Single Paxos state machine on top of each tablet ⇒ **Replication**
- Set of replicas is called **Paxos group**

- Every replica which is a leader implements
  - a **lock table** for concurrency control: state for 2PL
  - a **transaction manager** (TM) for distributed transactions
- Transaction manager used to implement **participant leader**
- If multiple Paxos groups are involved: Two-Phase Commit
  - One participant group is chosen as **coordinator group**
  - Participant leader of this group: **coordinator leader**
- State of each transaction manager is stored in Paxos group

# Spanserver Software Stack

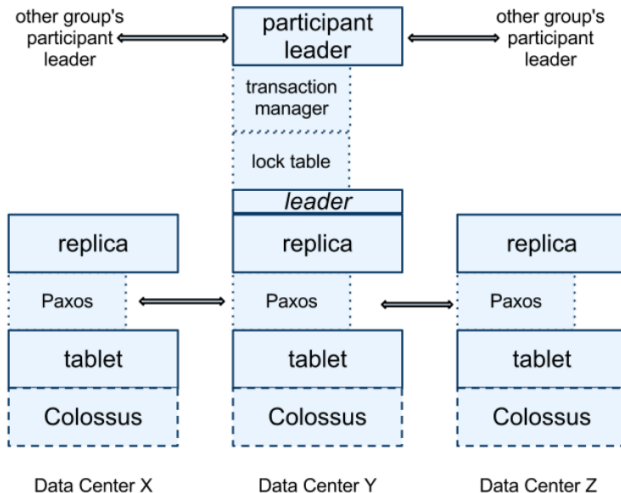


Figure 2 : Spanserver software stack

- **Directories:**

- Bucketing abstraction on top of tablet
- Contain set of contiguous keys sharing common prefix
- Control locality of data (physical)
- Unit of data placement (same replication configuration)
- Smallest unit to specify geographic-replication properties (**placement**)

- **Directories:**

- Bucketing abstraction on top of tablet
  - Contain set of contiguous keys sharing common prefix
  - Control locality of data (physical)
  - Unit of data placement (same replication configuration)
  - Smallest unit to specify geographic-replication properties (**placement**)
- Data is moved directory-wise between Paxos groups
  - Tablets may have multiple directories  
⇒ co-locate data frequently accessed together
  - Placements can be controlled in two dimensions:
    - ① Number & types of replicas
    - ② Geographic placement of replicas



# Directories & Placement




Figure 3 : Directories (the unit of data movement)

- Layered on top of directory-bucketed key-value mappings
- Applications create **databases**
- Database may contain any number of schematized **tables**
- Table contains rows, columns & versioned values

# Data Model

- Databases must be partitioned in one or more hierarchies
- **Directory table**: table at the top
- Directory contains hierarchically ordered rows
- Interleaving tables  $\Rightarrow$  locality relationships of multiple tables  
 $\Rightarrow$  significant for performance




Figure 4 : Interleaving example

Method	Returns
$TT.now()$	$TTinterval: [earliest, latest]$
$TT.after(t)$	true if $t$ has definitely passed
$TT.before(t)$	true if $t$ has definitely not arrived

Time is represented as a  $TTinterval$  (interval with bounded uncertainty)

**Guarantee:**  $TT.now().earliest \leq t_{abs}(e) \leq TT.now().latest$

- Time references: GPS and atomic clocks
- Set of **time master** machines per datacenter
- **Timeslave daemon** on each machine
- Masters compare time references regularly
- Daemons poll a variety of masters
- $\epsilon$  is the **instantaneous error bound** (typically 1 - 7ms)
- $\bar{\epsilon}$  is the average error bound (typically 4ms)
- $\epsilon$  is derived from applied worst-case local clock drift ( $200\mu s/sec$ )




Figure 5 : Visualization of  $TT.now()$  and  $\epsilon$

- TrueTime used to guarantee correctness properties to implement:
  - Externally-consistent transactions
  - Lock-free read-only transactions
  - Non-blocking reads in the past (snapshot reads)

# Timestamp Management

Operation	CC	Replica required
Read-Write transaction	Pessimistic	leader
Read-Only transaction	Lock-free	leader; any
Snapshot Read, client-provided timestamp	Lock-free	any
Snapshot Read, client-provided bound	Lock-free	any

- Standalone write  $\Rightarrow$  Read-Write (RW) transaction
- Non-snapshot standalone read  $\Rightarrow$  Read-Only (RO) transaction
- RO transactions must predeclare to not include writes
- RO transactions & snapshot reads proceed on any **sufficiently up-to-date** replica



- Timed leases for long-lived leadership (10s by default)
- Quorum of lease votes  $\Rightarrow$  leader has a lease
- Successful writes extend lease vote on replica
- **Disjointness invariant:**  
For each Paxos group, each Paxos leader's *lease interval* is disjoint from every other leader's
- Leaders must not abdicate before  $TT.after(s_{max})$
- $s_{max}$  . . . maximum timestamp used by a leader

- Uses 2PL: Timestamp assignment in between the two phases
- **Monotonicity invariant:**  
Within each Paxos group, Spanner assigns timestamps to Paxos writes in monotonically increasing order, even across leaders
- Enforced across leaders by disjointness invariant:  
Leader must only assign timestamps within interval of leader lease
- Timestamp  $s$  assigned  $\Rightarrow s_{max} = s$
- **External consistency invariant:**

$$t_{abs}(e_1^{commit}) < t_{abs}(e_2^{start}) \Rightarrow s_1 < s_2$$

# Read-Write Transactions

- Writes buffered at client until commit
- Client has completed all reads & buffered all writes  $\Rightarrow$  2PC
- Let  $e_i^{server}$  denote the arrival event of commit request for  $T_i$
- Two rules guarantee the external consistency invariant:

① **Start**

Coordinator leader assigns commit timestamp  
 $s_i \geq TT.now().latest$  after  $e_i^{server}$

② **Commit Wait**

Coordinator leader ensures clients cannot see  
data committed by  $T_i$  until  $TT.after(s_i)$  is true

# Read-Write Transactions

Acquire all locks




Figure 6 : RW transaction illustration

# Read-Write Transactions

Acquire all locks




Figure 6 : RW transaction illustration

# Read-Write Transactions

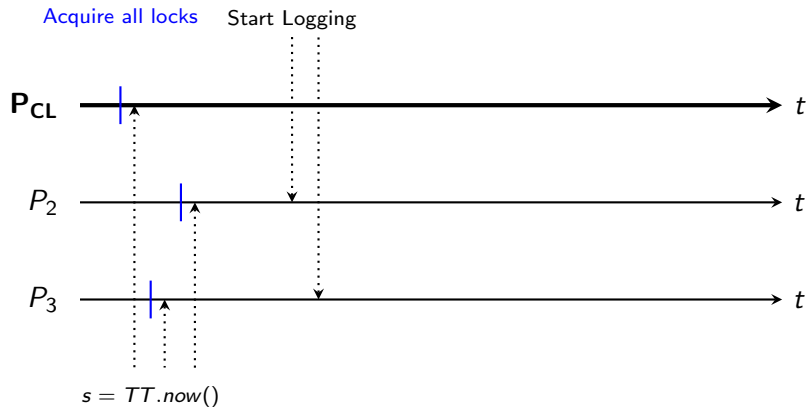


Figure 6 : RW transaction illustration

# Read-Write Transactions

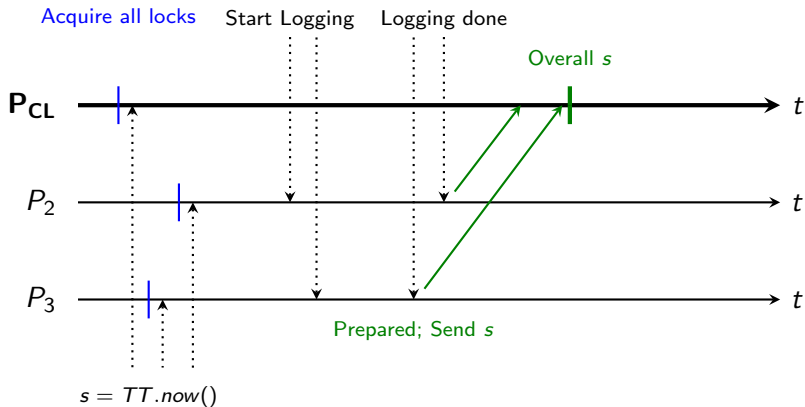


Figure 6 : RW transaction illustration

# Read-Write Transactions




Figure 6 : RW transaction illustration



# Read-Write Transactions

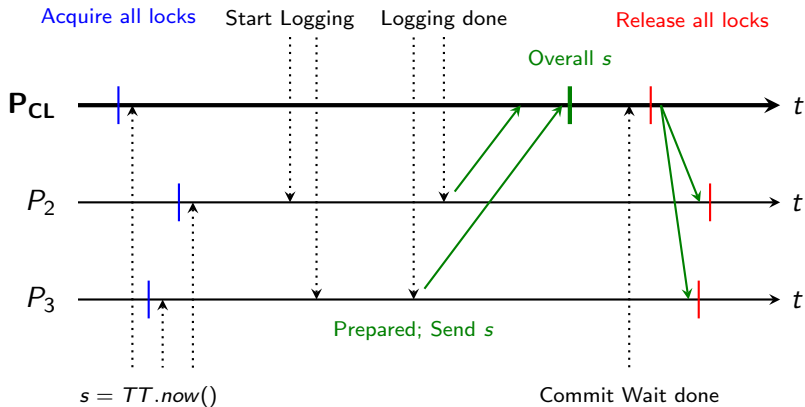


Figure 6 : RW transaction illustration

# Read-Write Transactions

Operation	CC	Replica required
Read-Write transaction	Pessimistic	leader

# Serving Reads at a Timestamp

Operation	CC	Replica required
Snapshot Read, client-provided timestamp	Lock-free	any
Snapshot Read, client-provided bound	Lock-free	any

- Monotonicity invariant allows to determine if a replica's state is **sufficiently up-to-date** to satisfy a read
- **Safe time**  $t_{safe}$  at each replica
- Replica satisfies read at timestamp  $t$  if  $t \leq t_{safe}$
- Performance benefits

# Read-Only Transactions

Operation	CC	Replica required
Read-Only transaction	Lock-free	leader; any

- Two phases:
  - 1 Assign timestamp  $s_{read}$
  - 2 Execute transaction's reads as snapshot reads at  $s_{read}$
- Spanner assigns  $s_{read}$  oldest timestamp to preserves external consistency

# Read-Only Transactions

- Negotiation phase between all Paxos groups involved
- **scope** expression summarizing the keys that will be read
- Scope's values served by single Paxos group
  - Client issues RO transaction to group leader
  - Leader assigns  $s_{read}$  = timestamp of last committed write
- Scope's values served by multiple Paxos groups
  - Negotiation between leaders possible
  - Simpler: reads execute at  $s_{read} = TT.now().latest$
- All reads can be sent to replicas that are sufficiently up-to-date

# Schema-Change Transactions

- Enabled by TrueTime
- Non-blocking variant of standard transaction
  - 1 Explicit assignment of a future timestamp
  - 2 Reads & writes synchronize with any registered schema-change timestamp at time  $t$
- Without TrueTime  $\Rightarrow$  schema change at  $t$  would be meaningless

- Setup:
  - 4GB RAM scheduling units per spanserver
  - 4 cores (AMD Barcelona 2200MHz) per spanserver
  - One spanserver per zone
  - Clients & zones in set of datacenters with network distance  $< 1\text{ms}$
  - Database created with 50 Paxos groups with 2500 directories
- Operations: standalone reads & writes of 4KB
- One unmeasured round of reads to warm any location caches

## Latency [ms] (less is better)

Replicas	Write	RO Transaction	Snapshot Read
1D	$9.4 \pm 0.6$	-	-
1	$14.4 \pm 1.0$	$1.4 \pm 0.1$	$1.3 \pm 0.1$
3	$13.9 \pm 0.6$	$1.3 \pm 0.1$	$1.2 \pm 0.1$
5	$14.4 \pm 0.4$	$1.4 \pm 0.05$	$1.3 \pm 0.04$

## Throughput [Kops/sec] (more is better)

Replicas	Write	RO Transaction	Snapshot Read
1D	$4.0 \pm 0.3$	-	-
1	$4.1 \pm 0.05$	$10.9 \pm 0.4$	$13.5 \pm 0.1$
3	$2.2 \pm 0.5$	$13.8 \pm 3.2$	$38.5 \pm 0.3$
5	$2.8 \pm 0.3$	$25.3 \pm 5.2$	$50.0 \pm 1.1$



# Availability

- Setup:

- Universe with 5 zones  $Z_i$ , each of which has 25 span servers
- All leaders explicitly placed in  $Z_1$
- 5s into each test, all servers in one zone were killed




Figure 6 : Effect of killing servers on throughput

- ① Is  $\epsilon$  truly a bound on clock uncertainty?
  - Local clock drifts  $> 200\mu\text{s}/\text{sec}$  would break assumptions
  - Clock issues infrequent relative to more serious hardware problems  
 $\Rightarrow$  TrueTime as trustworthy as any other piece of software

- ① Is  $\epsilon$  truly a bound on clock uncertainty?
  - Local clock drifts  $> 200\mu\text{s}/\text{sec}$  would break assumptions
  - Clock issues infrequent relative to more serious hardware problems  
 $\Rightarrow$  TrueTime as trustworthy as any other piece of software
- ② How bad does  $\epsilon$  get?




Figure 7 : Distribution of TrueTime  $\epsilon$  values

# Conclusions & Future Work

- 5 years from Spanner's inception to iterate to the current design
- Linchpin: TrueTime
- Build distributed systems with much stronger time semantics
- No longer depend on loosely synchronized clocks & weak time APIs

# Conclusions & Future Work

- 5 years from Spanner's inception to iterate to the current design
- Linchpin: TrueTime
- Build distributed systems with much stronger time semantics
- No longer depend on loosely synchronized clocks & weak time APIs
- Future Work:
  - Support parallel reads
  - Support direct changes of Paxos configurations
  - Improve single-node performance for complex queries
  - Support movement of processes between datacenters

Thank you for your attention!

Questions?