

# Spanner: Google's Globally-Distributed Database

Kocher Daniel

Summarizing Talk  
Efficient Algorithms Seminar  
University of Salzburg

January 27, 2016

# Outline

- 1 Motivation
- 2 Implementation
- 3 TrueTime
- 4 Concurrency Control
- 5 Evaluation
- 6 Conclusions & Future Work

# Motivation

- Distributed data at global scale
- Externally-consistent distributed transactions
- High availability
- Two "competitors": *Bigtable* & *Megastore*

- Distributed data at global scale
- Externally-consistent distributed transactions
- High availability
- Two "competitors": *Bigtable* & *Megastore*
- Why not use Bigtable?
  - Difficult to use for applications with complex, evolving schemas
  - Only **eventually-consistent** (no strong consistency)

# Motivation

- Distributed data at global scale
- Externally-consistent distributed transactions
- High availability
- Two "competitors": *Bigtable* & *Megastore*
- Why not use Bigtable?
  - Difficult to use for applications with complex, evolving schemas
  - Only **eventually-consistent** (no strong consistency)
- Why not use Megastore
  - Poor write throughput

**Spanner** evolved from a Bigtable-like versioned key-value store to a **temporal multi-version database**

**Spanner** evolved from a Bigtable-like versioned key-value store to a **temporal multi-version database**

- Versioned data is stored in schematized semi-relational tables
- Each version is automatically timestamped with its commit time
- Garbage collection & reads at old timestamps
- **General-purpose transactions**
- SQL-like query language

- Transaction  $i$  is denoted  $T_i$
- **Externally-consistent:**  
If  $T_1$  precedes  $T_2$  (Commit of  $T_1$  happens before Begin of  $T_2$ , no overlap), then  $T_1$  is serialized first



- Transaction  $i$  is denoted  $T_i$
- **Externally-consistent:**  
If  $T_1$  preceeds  $T_2$  (Commit of  $T_1$  happens before Begin of  $T_2$ , no overlap), then  $T_1$  is serialized first
- **Two-Phase Locking (2PL):**  
Guarantees serializability
  - 1 *Expanding* phase: locks are aquired, no locks are released
  - 2 *Shrinking* phase: locks are released, no locks are aquired
- **Two-Phase Commit (2PC):**  
Coordinates processes participating in atomic distributed transaction
  - 1 *Commit-Request* phase: Request "Yes" (Commit) or "No" (Abort) from every transaction process
  - 2 *Commit* phase: Commit transaction if all voted "Yes", otherwise abort

As a globally-distributed database, Spanner provides interesting features:

- Dynamically controlled replication configurations at fine grain
- Externally consistent reads & writes
- globally-consistent reads at a timestamp

⇒ enables **atomic schema changes** in presence of ongoing transactions

As a globally-distributed database, Spanner provides interesting features:

- Dynamically controlled replication configurations at fine grain
- Externally consistent reads & writes
- globally-consistent reads at a timestamp

⇒ enables **atomic schema changes** in presence of ongoing transactions

- Timestamps reflect serialization order (linearizability)
- Key enabler: novel TrueTime API exposing clock uncertainty

# Implementation (Top-down)

- Spanner deployment is called **universe**
- Organized as a set of **zones**
- Zones are the unit of administrative deployment & physical isolation
- Set of zones is the set of locations for replication

# Implementation (Top-down)

- Spanner deployment is called **universe**
- Organized as a set of **zones**
- Zones are the unit of administrative deployment & physical isolation
- Set of zones is the set of locations for replication
- Each zone consists of
  - several **spanservers** to store the data
  - a **zonemaster** to assign data to spanservers
  - a **location proxy** for clients to locate the assigned spanservers

# Implementation (Top-down)

- Spanner deployment is called **universe**
- Organized as a set of **zones**
- Zones are the unit of administrative deployment & physical isolation
- Set of zones is the set of locations for replication
- Each zone consists of
  - several **spanservers** to store the data
  - a **zonemaster** to assign data to spanservers
  - a **location proxy** for clients to locate the assigned spanservers
- **universe master**: console to display status information about zones
- **placement driver**:
  - handles automated data movement across zones
  - moves data due to updated replication constraints or load balancing

# Implementation (Top-down)




Figure 1 : Spanner server organization

- Responsible for 100 - 1000 **tablets**
- A tablet is a bag of key-value mappings  
(key:string, timestamp:int64) → string
- Timestamps are assigned to data ⇒ **multi-version database**
- Tablet states are stored on Colossus
- Single Paxos state machine on top of each tablet ⇒ **Replication**
- Set of replicas is called **Paxos group**



- Every replica which is a leader implements
  - a **lock table** for concurrency control: (key range) → lock state
  - a **transaction manager** (TM) for distributed transactions
- Transaction manager used to implement **participant leader**
- If multiple Paxos groups are involved: Two-Phase Commit
  - One participant group is chosen as **coordinator**
  - Participant leader of this group: **coordinator leader**
- State of each transaction manager is stored in Paxos group

# Spanserver Software Stack




Figure 2 : Spanserver software stack

- **Directories:**

- Bucketing abstraction on top of key-value mappings
- Contain set of contiguous keys sharing common prefix
- Control locality of data
- Unit of data placement (same replication configuration)
- Smallest unit to specify geographic-replication properties (**placement**)

- **Directories:**

- Bucketing abstraction on top of key-value mappings
  - Contain set of contiguous keys sharing common prefix
  - Control locality of data
  - Unit of data placement (same replication configuration)
  - Smallest unit to specify geographic-replication properties (**placement**)
- Data is moved directory-wise between Paxos groups
  - Tablets may have multiple **partitions** of row space  
⇒ co-locate directories frequently accessed together

- **Directories:**

- Bucketing abstraction on top of key-value mappings
- Contain set of contiguous keys sharing common prefix
- Control locality of data
- Unit of data placement (same replication configuration)
- Smallest unit to specify geographic-replication properties (**placement**)
- Data is moved directory-wise between Paxos groups
- Tablets may have multiple **partitions** of row space  
⇒ co-locate directories frequently accessed together
- Placements can be controlled in two dimensions:
  - ① Number & types of replicas
  - ② Geographic placement of replicas
- **Movedir:** background task to move directories

# Directories & Placement




Figure 3 : Directories (the unit of data movement)

- Features:
  - Schematized semi-relational tables
  - SQL-like query language
  - General-purpose transactions
- Layered on top of directory-bucketed key-value mappings
- Applications create **databases** in universe
- Database may contain unlimited number of schematized **tables**
- Table contains rows (must have names), columns & versioned values

# Data Model

- Databases must be partitioned in one or more hierarchies
- **Directory table**: table at the top
- Directory contains ordered rows starting at a directory table key
- Interleaving tables to form directories is significant
- Locality relationships between multiple tables  $\Rightarrow$  good performance




Figure 4 : Interleaving example



Method	Returns
$TT.now()$	$TTinterval$ : $[earliest, latest]$
$TT.after(t)$	true if $t$ has definitely passed
$TT.before(t)$	true if $t$ has definitely not arrived

Table 1 : TrueTime API ( $t$  is of type  $TTstamp$ )

Time is represented as a  $TTinterval$  (interval with bounded uncertainty)

Let  $e$  be an event and  $t_{abs}(e)$  denote the absolute time of  $e$ .  
For  $tt = TT.now()$ , the following property holds:

$$tt.earliest \leq t_{abs}(e) \leq tt.latest$$

- Time references: GPS and atomic clocks
- Set of **time master** machines per datacenter
- **Armageddon masters**: masters using atomic clocks
- **Timeslave daemon** on each machine
- Masters compare their time references regularly
- Daemons poll a variety of masters (liar detection)




Figure 5 : Visualization of  $TT.now()$  and  $\epsilon$

# Concurrency Control

- TrueTime is used to guarantee correctness properties
- Properties are used to implement
  - Externally-consistent transactions
  - Lock-free read-only transactions
  - Non-blocking reads in the past

# Timestamp Management

Operation	CC	Replica required
Read-Write transaction	Pessimistic	leader
Read-Only transaction	Lock-free	leader; any
Snapshot Read, client-provided timestamp	Lock-free	any
Snapshot Read, client-provided bound	Lock-free	any

Table 1 : Supported types of reads & writes

- Standalone write  $\Rightarrow$  RW transaction
- Non-snapshot standalone read  $\Rightarrow$  RO transaction
- RO transactions
  - must predeclare to not include writes
  - can proceed on any sufficiently up-to-date replica (also snapshot reads)

- Timed leases for long-lived leadership (10s by default)
- Quorum of lease votes  $\Rightarrow$  leader has a lease
- Successful writes extend lease vote on replica
- **Disjointness invariant:**  
For each Paxos group, each Paxos leader's *lease interval* is disjoint from every other leader's
- Leaders must not abdicate before  $TT.after(s_{max})$ , where  $s_{max}$  denotes the maximum timestamp used by a leader

# Timestamps in RW Transactions

- Timestamp assignment only in between the two phases (2PL)
- **Monotonicity variant:**  
Within each Paxos group, Spanner assigns timestamps to Paxos writes in monotonically increasing order, even across leaders
- Enforced across leaders by disjointness invariant:  
Leader must only assign timestamps within the interval of leader lease
- Timestamp  $s$  assigned  $\Rightarrow s_{max} = s$
- **External consistency invariant:**

$$t_{abs}(e_1^{commit}) < t_{abs}(e_2^{start}) \Rightarrow s_1 < s_2$$

# Timestamps in RW Transactions

- Let  $e_i^{server}$  denote the arrival event of commit request for  $T_i$
- Two rules guarantee the external consistency invariant:
  - 1 **Start**  
Coordinator leader assigns commit timestamp  
 $s_i \geq TT.now().latest$  after  $e_i^{server}$
  - 2 **Commit Wait**  
Coordinator leader ensures clients cannot see  
data committed by  $T_i$  until  $TT.after(s_i)$  is true



# Serving Reads at a Timestamp

- Monotonicity invariant allows to determine if a replica's state is **sufficiently up-to-date** to satisfy a read
- **Safe time**  $t_{safe}$  at each replica
- Replica satisfies read if  $t \leq t_{safe}$
- $t_{safe} = \min(t_{safe}^{Paxos}, t_{safe}^{TM})$

# Timestamps in RO Transactions

- Two phases:
  - 1 Assign timestamp  $s_{read}$
  - 2 Execute transaction's reads as snapshot reads at  $s_{read}$
- Spanner assigns the oldest timestamp that preserves external consistency to reduce the changes of blocking

# Read-Write Transactions

- RW transactions buffered at client until Commit
- Reads do not see effects of transaction's writes
- Client has completed all reads & buffered all writes  $\Rightarrow$  2PC

# Read-Write Transactions

- RW transactions buffered at client until Commit
- Reads do not see effects of transaction's writes
- Client has completed all reads & buffered all writes  $\Rightarrow$  2PC
- Non-coordinator-participant leader
  - 1 Chooses prepare timestamp according to monotonicity
  - 2 Logs prepare record through Paxos
  - 3 Notifies coordinator of its prepare timestamp

# Read-Write Transactions

- RW transactions buffered at client until Commit
- Reads do not see effects of transaction's writes
- Client has completed all reads & buffered all writes  $\Rightarrow$  2PC
- Non-coordinator-participant leader
  - 1 Chooses prepare timestamp according to monotonicity
  - 2 Logs prepare record through Paxos
  - 3 Notifies coordinator of its prepare timestamp
- Coordinator leader
  - 1 Chooses timestamp for whole transaction
  - 2 Commit timestamp  $s$  must be
    - $\geq$  all prepare timestamps
    - $> TT.now().latest$  at the time the Commits were received
    - $>$  any previously assigned timestamp
  - 3 Logs Commit record through Paxos

# Read-Write Transactions

Acquire all locks




Figure 6 : RW transaction illustration

# Read-Write Transactions

Acquire all locks




Figure 6 : RW transaction illustration

# Read-Write Transactions

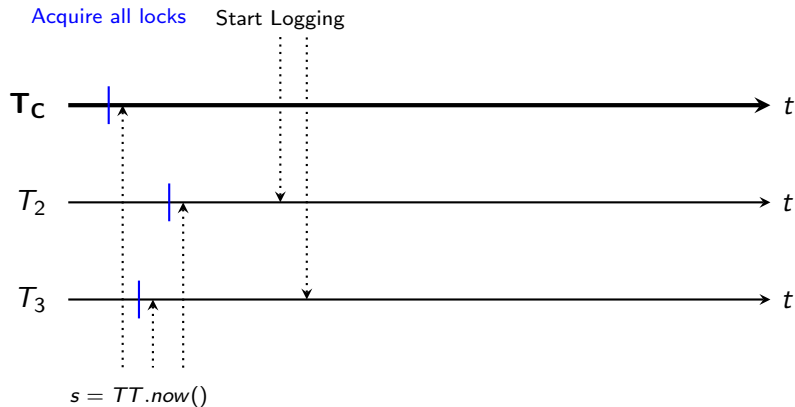


Figure 6 : RW transaction illustration



# Read-Write Transactions

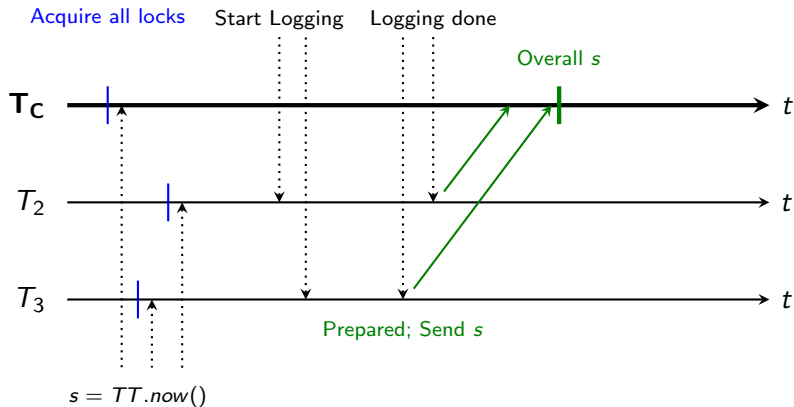


Figure 6 : RW transaction illustration

# Read-Write Transactions




Figure 6 : RW transaction illustration

# Read-Write Transactions




Figure 6 : RW transaction illustration

# Read-Only Transactions

- **scope** expression summarizing the keys that will be read
- Served by single Paxos group  
⇒ client issues transaction to group leader
- Served by multiple Paxos groups  
⇒ client executes reads at  $s_{read} = TT.now().latest$
- All reads can be sent to replicas that are sufficiently up-to-date

# Schema-Change Transactions

- TrueTime enables this feature
- Use of standard transaction infeasible
- Non-blocking variant of standard transaction
  - ① Explicit assignment of a future timestamp (Prepare phase)
  - ② Reads & writes synchronize with any registered schema-change timestamp at time  $t$
- Without TrueTime  $\Rightarrow$  schema change at  $t$  would be meaningless

- Setup:
  - 4GB RAM scheduling units per spanserver
  - 4 cores (AMD Barcelona 2200MHz) per spanserver
  - One spanserver per zone
  - Clients & zones in set of datacenters with network distance  $< 1\text{ms}$
  - Database created with 50 Paxos groups with 2500 directories
- Operations: standalone reads & writes of 4KB
- One unmeasured round of reads to warm any location caches

## Latency [ms] (less is better)

Replicas	Write	RO Transaction	Snapshot Read
1D	$9.4 \pm 0.6$	-	-
1	$14.4 \pm 1.0$	$1.4 \pm 0.1$	$1.3 \pm 0.1$
3	$13.9 \pm 0.6$	$1.3 \pm 0.1$	$1.2 \pm 0.1$
5	$14.4 \pm 0.4$	$1.4 \pm 0.05$	$1.3 \pm 0.04$

Table 2 : Latency experiments

- Sufficiently few operations to avoid queuing
- Increasing number of replicas:
  - Latency stays roughly constant with less standard deviation
  - Latency to achieve quorum less sensitive to stragglers

## Throughput [Kops/sec] (more is better)

Replicas	Write	RO Transaction	Snapshot Read
1D	$4.0 \pm 0.3$	-	-
1	$4.1 \pm 0.05$	$10.9 \pm 0.4$	$13.5 \pm 0.1$
3	$2.2 \pm 0.5$	$13.8 \pm 3.2$	$38.5 \pm 0.3$
5	$2.8 \pm 0.3$	$25.3 \pm 5.2$	$50.0 \pm 1.1$

Table 2 : Throughput experiments

- Sufficiently many operations to saturate servers' CPUs
- Increasing number of replicas:
  - Snapshot read throughput increases almost linear



- Setup:
  - Universe with 5 zones  $Z_i$ , each of which has 25 spanservers
  - Database sharded into 1250 Paxos groups
  - 100 clients constantly issued snapshot reads (50K reads/sec)
  - All leaders explicitly placed in  $Z_1$
  - 5s into each test, all servers in one zone were killed




Figure 7 : Effect of killing servers on throughput

- Availability benefits of running Spanner in multiple datacenters

- Two questions:
  - Is  $\epsilon$  truly a bound on clock uncertainty?
  - How bad does  $\epsilon$  get?

- Two questions:
  - Is  $\epsilon$  truly a bound on clock uncertainty?
  - How bad does  $\epsilon$  get?
- Answers:
  - Clock issues infrequent relative to more serious hardware problems  
 $\Rightarrow$  TrueTime as trustworthy as any other piece of software

# TrueTime

- Two questions:
  - Is  $\epsilon$  truly a bound on clock uncertainty?
  - How bad does  $\epsilon$  get?
- Answers:
  - Clock issues infrequent relative to more serious hardware problems  
⇒ TrueTime as trustworthy as any other piece of software




Figure 8 : Distribution of TrueTime  $\epsilon$  values

# Conclusions & Future Work

- 5 years from Spanner's inception to iterate to the current design
- Combines & extends on ideas from database & systems community
- Linchpin: TrueTime
- Build distributed systems with much stronger time semantics
- No longer depend on loosely synchronized clocks & weak time APIs

# Conclusions & Future Work

- 5 years from Spanner's inception to iterate to the current design
- Combines & extends on ideas from database & systems community
- Linchpin: TrueTime
- Build distributed systems with much stronger time semantics
- No longer depend on loosely synchronized clocks & weak time APIs
- Future Work:
  - Implement optimistical reads in parallel
  - Support direct changes of Paxos configurations
  - Improve single-node performance
  - Support movement of processes between datacenters

Thank you for your attention!

Questions?